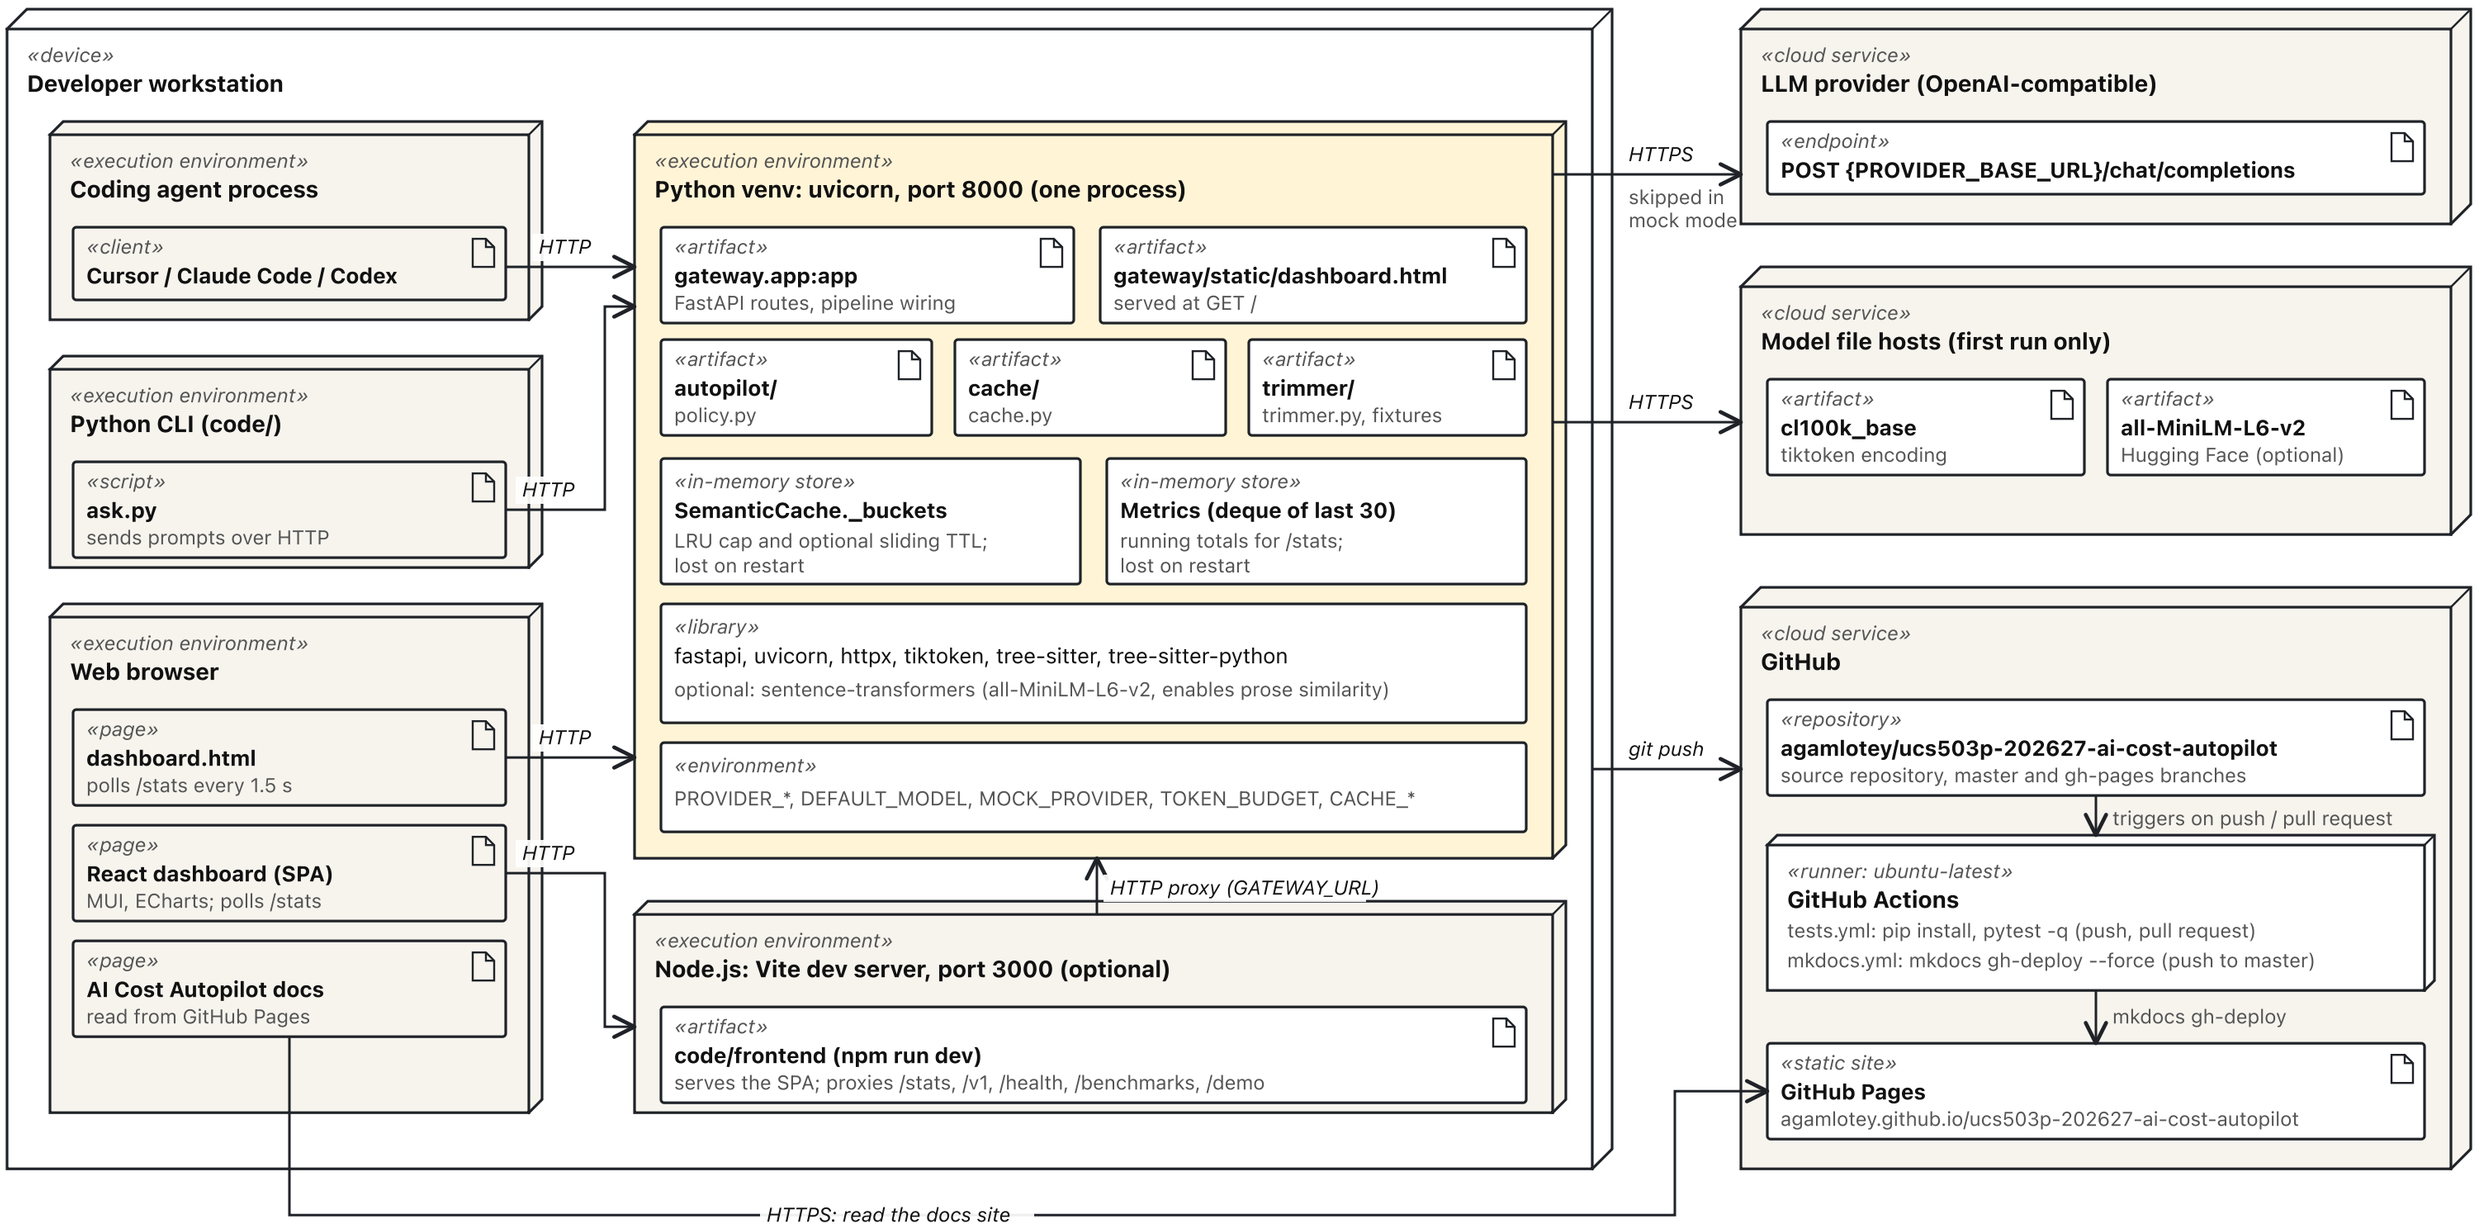
Deployment Diagram — AI Cost Autopilot

What the repository actually deploys today: one local gateway process per developer. There is no container image or hosted production gateway.



Notes

- Mock mode: with MOCK_PROVIDER=1 or no PROVIDER_API_KEY, the gateway answers locally and never contacts the LLM provider, so the whole pipeline runs offline at no cost.
- Base URL: a coding agent connects by pointing its base URL at http://localhost:8000/v1. The built-in dashboard needs no separate server: GET / serves dashboard.html from the same process.
- State: the cache and metrics live inside the single uvicorn process. They are not shared between workers and do not survive a restart. A persistent or shared store is future work.
- Environment: the code reads settings with os.getenv and does not load .env by itself; export them or pass uvicorn --env-file .env.
- Offline fallbacks: with no tiktoken download, count_tokens falls back to len/4. Without sentence-transformers, the cache still matches exact text and only prose similarity is off.
- demo.py runs in process through FastAPI TestClient with the mock provider, so it needs no running server and is not a separate deployment.
- Not deployed: docs/architecture.md lists SQLite/FAISS as the vector store, but the code does not use them. No Dockerfile, compose file or hosted gateway exists in the repository.

Notation node (device, environment or service) · deployed artifact or store · communication path, labelled with its protocol; the arrow points at the callee