

Ai **COST AUTOPILOT**

AN AGENTIC GATEWAY FOR TOKEN OPTIMIZATION

same brain, smaller bills

presented by: furmaan, agam, devansh

The INTRO

- AI coding tools (Cursor, Claude Code, Codex) are now part of everyday development.
- Every request sends code to an LLM — and you pay for every token.
- Our project: a proxy that sits between the coding agent and the AI, cutting cost and code exposure automatically.
- Goal: up to ~40% lower AI bills, with no change to how developers work.



OpenAI



Claude




CURSOR



Gemini

The PROBLEM

- Coding agents send large chunks of the codebase on every request.
- You pay per token → bills grow unpredictably as teams and code grow.
- Prompts are often bloated and repetitive — paying twice for the same thing.
- Privacy risk: private code is shipped to the cloud in bulk.
- No standard layer exists today to control this spend.



why now **WHY THIS MATTERS**

THE PROBLEM IS EXPLODING

- AI coding adoption is skyrocketing across every dev team.
- LLM API costs are becoming a top, unpredictable line item in engineering budgets.

TODAY'S "SOLUTIONS" FALL SHORT

- Generic text compressors cut blindly — risky, and they can break code.
- Teams are left choosing between cost, quality, and privacy — pick two.

THE OPPORTUNITY

- Most tokens spent are avoidable: repeated questions and unneeded code.
- A smart layer can cut waste and reduce how much code leaves the machine — cost and privacy solved together



**Compiler-Aware
Trimmer**



**Semantic
Cache**



Autopilot



**Universal
Plugin**



Private by design

SOLUTION

Cut your AI bills by up to 40% — without breaking your work. A smart middleman that sits between you and the AI, saving money on every request, automatically.

THANK
You