

REQUEST FLOW



CODING AGENT

(Cursor / Claude Code / Codex)

POST /v1/chat/completions



GATEWAY (FastAPI)

1

snapshot the original request

2

compute cheap signals (token count, has_code?)

3

autopilot.decide(request, signals) -> plan

4

if plan.use_cache -> cache.lookup()

hit



HIT
RETURN

5

if plan.trim -> trimmer.trim()

6

provider.forward() -> real LLM (or mock)



LLM provider

7

cache.store(original, response)

8

return response



RESPONSE

returned to coding agent